

# Decentralized Recommendation System

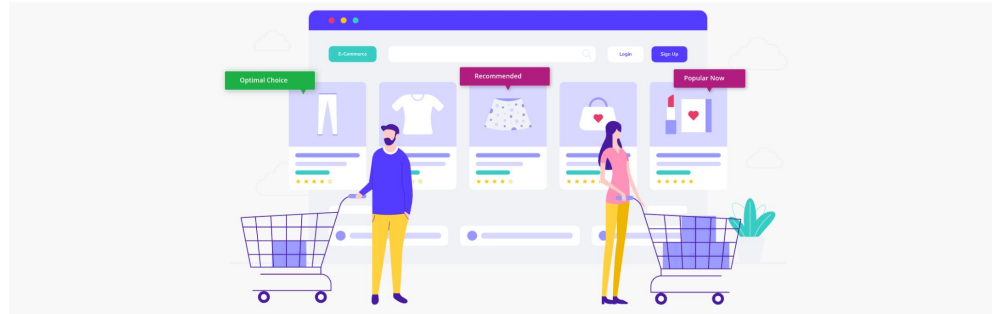
Matthew Kuo, Laura Li, Vivian Xiao, Megan Yang

April 22, 2025

# Problem Definition

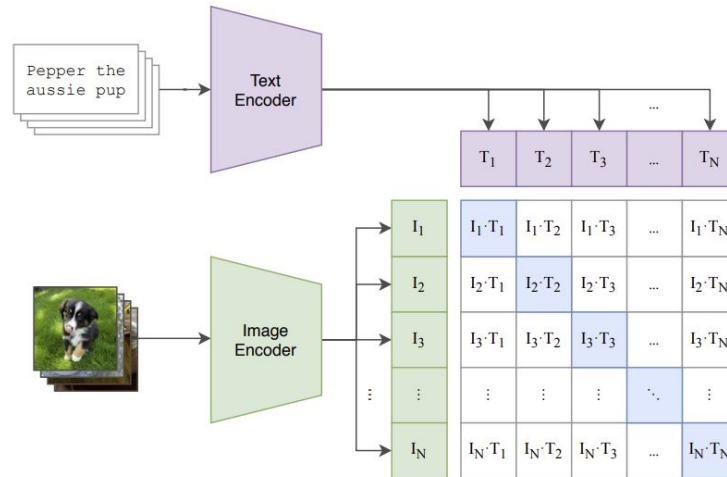
# Problem Statement

- **Problem:**
  - Current systems are platform-specific and disconnected
  - Poor performance in cold-start scenarios
    - Users have no interaction history
- **Goal:**
  - Cross-platform recommendations that leverages product metadata
  - Personalization for new users



# Key Terms

- **CLIP**: A pre-trained model that encodes images and texts
- **User embeddings**: Vectors of a users' preferences based on ratings
- **Item embeddings**: Vectors of the visual/textual features generated by CLIP
- **Synthetic Data**: Artificially generated data used to simulate the real-world for training/testing



# Data Preparation

# Dataset 1: Images + Captions

- ~**3000** data points (clothing items with metadata)
  - ~1500 men clothes from Myntra and ~1500 women clothes from ASOS

name	description	price
Mid-Rise Wide-Leg Cargo Pants	A pair of twill pants featuring a mid-rise waist, belt loops, zip fly and button-front closure, slanted front pockets, wide leg, leg cargo flap pockets with frayed trim, and back patch pockets.	24.49



	text_embedding	image_embedding
0	[0.034454345703125, 0.4833984375, -0.090270996...	[0.1217041015625, 0.1280517578125, -0.25146484...
1	[0.08929443359375, 0.05108642578125, -0.151855...	[0.07623291015625, 0.62255859375, -0.115661621...

## Dataset 2: Personas

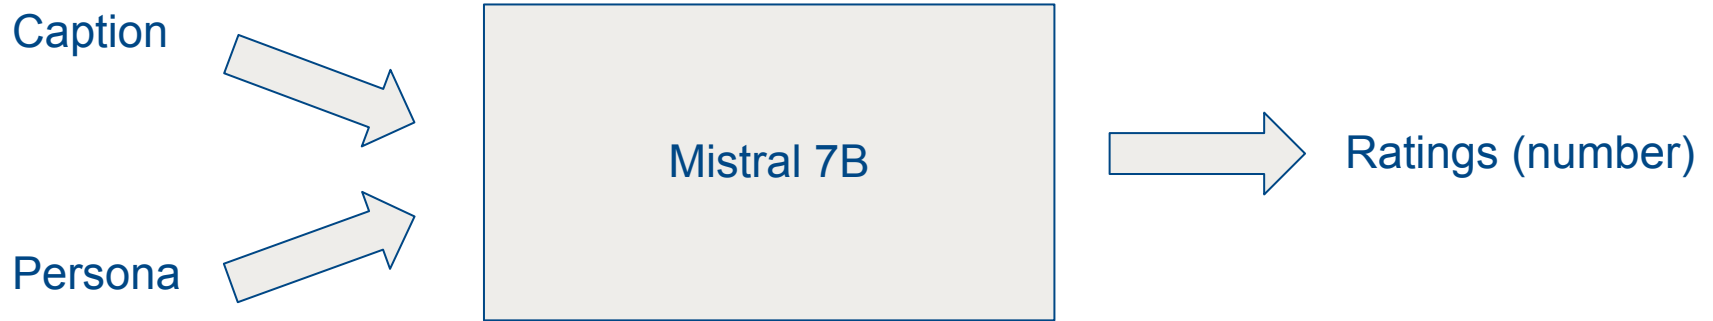
- Synthetic people with varying opinions on what they like
- 30 personas
- Generated by ChatGPT

- Example:

```
{"name": "Alex", "bio": "A 28-year-old graphic designer favoring Scandinavian minimalist styles. Prefers monochrome palettes (black, white, grey), high-quality natural fabrics (linen, wool), clean geometric cuts, and avoids logos or excessive detailing."}
```

## Dataset 3: Ratings

- Matrix (80% sparsity) of what each person thinks about each item





# Technical Approach

# Model Landscape Overview

Model	Uses Metadata (Image/Text embeddings)	Uses User Ratings (Synthetic Data)	Scalability (# of users/items)	Recommendation Type
Content Filtering	✓	✗	Easy (per user basis)	Uniform but personalized
Collaborative Filtering	✗	✓	Challenging (pairwise similarities)	Novel, social-based
Low-Rank Completion	Optional	✓	Moderate (high initial cost)	Interpolative
Two-Tower	✓	✓	Moderate (high initial cost)	Hybrid, rich representations

# Literature Review

## 1. Collaborative Filtering

- Amazon's item-to-item collaborative filtering
- Linden, G., Smith, B., & York, J. (2003). Amazon.com Recommendations: Item-to-Item Collaborative Filtering. *IEEE Internet Computing*, 7(1), 76–80.

## 2. Content-Based Filtering

- Spotify's content-based recommendation system
- Bangera, S., Nagaonkar, V., Tiwari, A., Ansari, S., & Talekar, K. (2024). Spotify Recommendation System. *International Research Journal of Modernization in Engineering, Technology and Science*, 6(2).

## 3. Low-Rank Matrix Completion


- Netflix's matrix factorization approach
- Amatriain, X., & Basilico, J. (2015). Recommender Systems in Industry: A Netflix Case Study. In *Recommender Systems Handbook* (pp. 385–419). Springer.

## 4. Two-Tower Neural Networks

- YouTube's deep neural networks for recommendations
- Covington, P., Adams, J., & Sargin, E. (2016). Deep Neural Networks for YouTube Recommendations. In *Proceedings of the 10th ACM Conference on Recommender Systems* (pp. 191–198). ACM.

# Collaborative Filtering

- **User-based:** finding similar users and suggesting what they like



## Implementation Steps

1. Extract user/item vectors from dataset.
2. Compute cosine similarity between the target user and others.
3. Select top-N similar users/items as weights.
4. Predict item scores using weighted preferences.
5. Rank & recommend top items based on scores.

## Math Equations

$$\hat{r}_{u,i} = \frac{\sum_{v \in N_u} \text{sim}(u,v) \cdot r_{v,i}}{\sum_{v \in N_u} |\text{sim}(u,v)|}$$

- $\hat{r}_{u,i}$ : predicted rating for user  $u$  on item  $i$
- $r_{v,i}$ : actual rating of user  $v$  on item  $i$
- $\text{sim}(u,v)$ : similarity (e.g., cosine) between user  $u$  and  $v$
- $N_u$ : top-N similar users to user  $u$

# Content-based Filtering

- Analyzes **item features** (e.g., descriptions, image embeddings) and compares them to a user's **past preferences**.
- **User preference vector** ( $v_u$ ) created by averaging feature representations of liked items ( $L_u$ ).
- New items ( $j$ ) ranked based on **cosine similarity** ( $s_j$ ) to the user preference vector.
  - Highest-scoring items recommended


## Mathematical Formulation

$$v_u = \frac{1}{|L_u|} \sum_{i \in L_u} x_i$$

$$\arg \max_{j \notin L_u} s_j \quad s_j = \frac{v_u \cdot x_j}{\|v_u\| \|x_j\|}$$

## Implementation Steps

1. Represent each item as a feature vector (text & image embeddings)
2. Compile a list of all “liked” items for a user
3. Calculate the cosine similarity between the preference vector and all other items.
4. Sort items by similarity.
5. Return top recommendations



# Low Rank Matrix Completion

We model the rating matrix  $R \in \mathbb{R}^{n_{\text{items}} \times n_{\text{users}}}$  as the product of two low-rank matrices:

$U \in \mathbb{R}^{n_{\text{users}} \times r}$  : user latent factors

$V \in \mathbb{R}^{n_{\text{items}} \times r}$  : item latent factors



Objective Function:

$$\min_{U, V} \sum_{(i, j) \in \text{observed}} (R_{ij} - \langle U_j, V_i \rangle)^2 + \lambda (\|U\|_F^2 + \|V\|_F^2)$$

Minimize reconstruction error only on observed entries

Goal: Fill out a partially observed user-item rating matrix using a low-rank factorization approach.

		-1		
			1	
1	1	-1	1	-1
1				-1
		-1		

1	1	-1	1	-1
1	1	-1	1	-1
1	1	-1	1	-1
1	1	-1	1	-1
1	1	-1	1	-1

Implementation Steps:

1. Convert the sparse rating matrix into training triplets (user\_id, item\_id, rating)
2. Initialize  $U$ ,  $V$  using PyTorch nn.Embedding
3. Predictions are computed as the dot product
4. Optimize with mini-batch gradient descent using MSE loss

Hyperparameter Tuning:

- Performed grid search over rank (2..30) and learning rate
- Selected best model based on Precision@10 on validation data

# Low Rank Variations

## Projection Layer with Item Embeddings

- Item embeddings (text+image) are projected into low-rank space via a fixed linear layer.
- User factors are learned
- Good for cold-start items since item embeddings are known upfront.

## Pairwise Ranking Loss

Bayesian Personalized Ranking (BPR) loss:

$$\mathcal{L}_{\text{BPR}} = - \sum_{(u,i,j)} \log \sigma(\langle U_u, V_i \rangle - \langle U_u, V_j \rangle)$$

- Optimizes pairwise ranking: push relevant items above irrelevant ones.
- Captures relative ranking positions

Model Type	Objective	Pros	Cons
Classic	Rating prediction	Simple, effective	No metadata support
Projection-based	Cold-start generalization	Leverages image/text features	May underfit latent needs
BPR (Pairwise)	Ranking optimization	Directly optimizes ranking	Harder to train/stabilize

# Two-Tower

Instead of learning one large joint representation of users and items

- Use one NN (tower) to learn user reps and the other to learn item reps.
- Compare them with a similarity function.

## Mathematical Formulation

1. Use CLIP to provide initial item embeddings.
2. Obtained input for the user tower:

$$e_i = \left[ \underbrace{\text{CLIP}_{\text{image}}(x_i)}_{512}, \underbrace{\text{CLIP}_{\text{text}}(x_i)}_{512} \right] \in \mathbb{R}^{1024}$$

$$s \in \mathbb{R}^N, \quad s_i = \begin{cases} +1, & i \in \mathcal{S}, \\ -1, & i \notin \mathcal{S}. \end{cases}$$

3. User embeddings are passed through user tower for transformation:

$$\text{a. } u = \text{ReLU}(W_u s + b_u), \quad W_u \in \mathbb{R}^{d \times N}, \quad b_u \in \mathbb{R}^d$$

4. Item embeddings are fed into item tower:

$$\text{b. } v_i = \text{ReLU}(W_v e_i + b_v), \quad W_v \in \mathbb{R}^{d \times d_e}, \quad b_v \in \mathbb{R}^d$$

5. Normalized the outputs with L2-norm.

6. Cosine similarity per user-item pair.

$$\text{sim} = \cos(\bar{u}, \bar{v}) \in [-1, 1]$$

7. Rating Predictor MLP that takes in sim and returns r

8. Loss function:

$$\mathcal{L} = \frac{1}{|\mathcal{B}|} \sum_{(u,i) \in \mathcal{B}} (r_i - y_i)^2,$$



Demo!

# Results

# Performance Metrics (Part 1)

## Root Mean Error Square

Measures the square root of the average squared difference between predicted and true ratings.

$$\text{RMSE} = \sqrt{\frac{1}{N} \sum_{(i,j) \in \text{Val}} (r_{ij} - \hat{r}_{ij})^2}$$

## Mean Absolute Error

Measures the average absolute difference between predicted and true ratings. Less sensitive to outliers than RMSE.

$$\text{MAE} = \frac{1}{N} \sum_{(i,j) \in \text{Val}} |r_{ij} - \hat{r}_{ij}|$$

## Performance Metrics (Part 2)

### Precision@10

Measures the fraction of top-10 recommended items that a user would buy.

$$\text{Precision@10} = \frac{\text{\#relevant items in top-10}}{10}$$

### Recall@10

Measures the fraction of all items a user would buy that appear in the top-10 recommendations.

$$\text{Recall@10} = \frac{\text{\#relevant items in top-10}}{\text{\#relevant items}}$$

# Model Metrics – Collaborative filtering

## Collaborative Filtering Evaluation (Original vs Binary)



### Full Ratings

RMSE: 3.2663

MAE: 2.6129

Precision: 0.0167

Recall: 0.0042



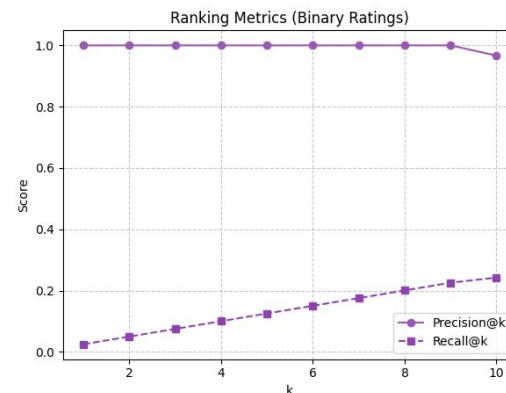

### Binary Ratings

RMSE: 1.5316

MAE: 1.1729

Precision: 0.9667

Recall: 0.2426



# Model Metrics – Content Based Filtering

## Content-Based Filtering Evaluation (Original vs Binary)

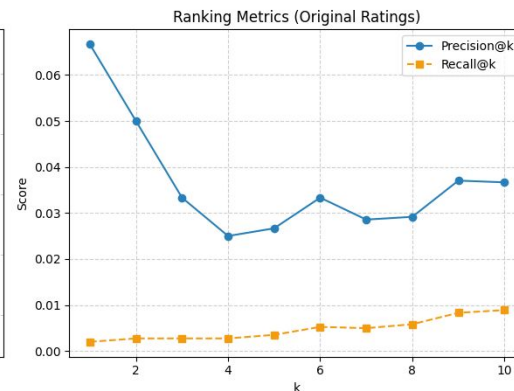

### Full Ratings

RMSE: 3.1535

MAE: 2.5099

Precision: 0.0367

Recall: 0.0093





### Binary Ratings

RMSE: 1.5474

MAE: 1.1972

Precision: 1.0000



Recall: 0.2510



# Model Metrics – Baseline Low Rank Model



## Full Ratings (with rank=22)

RMSE: 3.4647  
MAE: 2.8413  
Precision@10: 0.0367  
Recall@10: 0.0099



## Binary Ratings (with rank=27)

RMSE: 1.0942  
MAE: 1.0094  
Precision@10: 0.0333  
Recall@10: 0.0086




# Model Metrics – Low Rank Projection

## Full Ratings (with rank=20)

RMSE: 2.6131  
MAE: 2.1946  
Precision@10: 0.0133  
Recall@10: 0.0030

## Binary Ratings (with rank=16)

RMSE: 0.9809  
MAE: 0.9639  
Precision@10: 0.0133  
Recall@10: 0.0030





# Model Metrics – Low Rank with Pairwise Ranking Loss


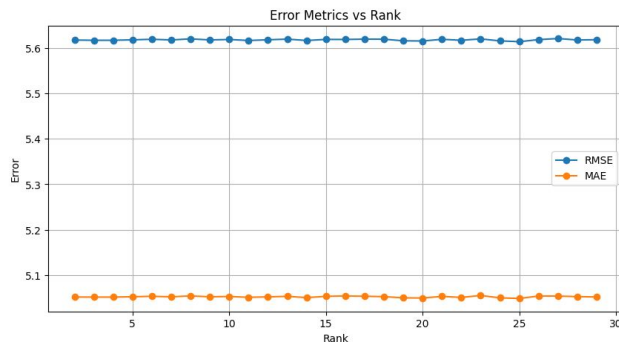
## Full Ratings (with rank=11)

RMSE: 5.6161

MAE: 5.0515

Precision: 0.0367

Recall: 0.0092





## Binary Ratings (with rank=5)

RMSE: 0.9998

MAE: 0.9989

Precision: 0.0333

Recall: 0.0084



# Model Metrics – Two Tower

## Full Ratings

RMSE: 2.8050

MAE: 2.2545

Precision@10: 0.0333

Recall@10: 0.0076



## Binary Ratings

RMSE: 1.1365

MAE: 0.9426


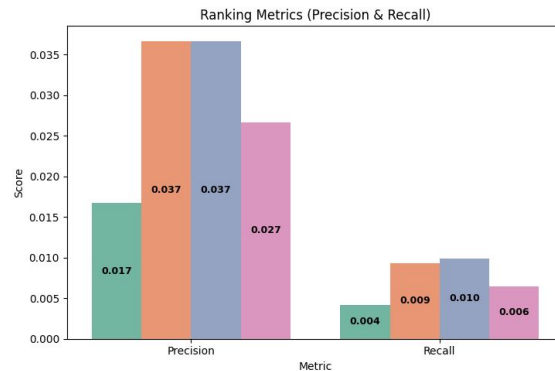
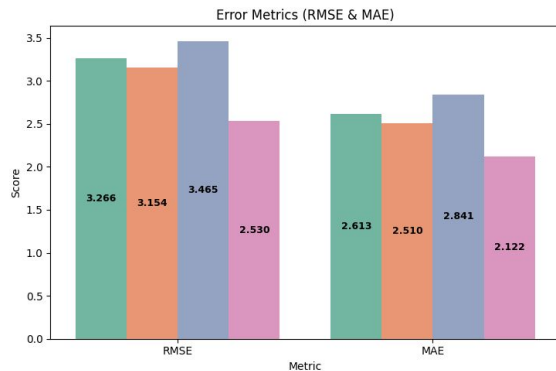
Precision: 0.0133

Recall: 0.0030







# Performance Summary

## Full Ratings



## Binary Ratings



Reflection

# Hardest Technical/Conceptual Difficulty

- Conceptually understanding each algorithm and how to measure their performance
  - Initial results for collaborative filtering and two towers were poor
  - Trial and Error for different LLMs (Mistral, Llama, Phi-2)
  - Tune hyperparameters and find the best precision@k
- Creating the necessary datasets
  - Web scraping was largely infeasible due to website security controls/resource constraints
  - Generating usable, synthetic personas and ratings

# Workflow

- **Easier:** Content-based filtering
  - Computed pairwise similarities and gave great results
- **Harder:** Collecting data
  - Planned to scrape images and recruit volunteers to “like” or “dislike” them
  - E-commerce sites blocked the scraping, and labeling was too manual
  - Limited online resources for clothing dataset with metadata and high quality images

# Evolving Goals

- **Initial:** Implement all four algorithms and compare outputs
- **Mid-project pivot:**
  - Create synthetic data and see how algorithms behave
    - Appends a new column of +1/-1 or NaNs so that algorithms can use the new user
  - Automated data-cleaning pipeline

# AI Tools Assist

- Persona & ratings generation
- Initial model training & debugging
- Model exploration
  - Variations of low-rank models (e.g., fixed projection, BPR)





# Individual Contributions

# Megan

- Most Surprising Result or Finding

- One of the best results were content filtering even though it was so simple.

- Specific lecture

- Lecture 9 helped us choose Adam over Adagrad because Adam retains Adagrad's per-parameter adaptive scaling—automatically dampening parameters with large gradients. Thus, our two-tower network reached useful recommendation quality in fewer epochs

- Perspective on optimization

- Thought optimization was simple and theoretical. In practice, however, nonconvex problems behave unpredictably and some practices are more practical although less optimal (for example step-size, we should be diminishing but choose a constant step-size and manually decrease it).

- 2 more weeks

- Set up a hyperparameter-optimization pipeline to explore learning rates, layer sizes, and regularization strengths for the two-tower

- Restart the project

- Prioritize data collection infrastructure first—designing a user-friendly labeling interface and recruitment plan—before implementing multiple algorithms.

# Vivian

- **Most Surprising Result or Finding**
  - Data quality and preprocessing ended up being as important for performance as model selection
- **Most useful lecture concept**
  - Problem Formulation in PyTorch: The focus on defining clear objectives and leveraging autograd for gradients made implementing new models in PyTorch easier.
- **Perspective change**
  - Appreciate the trade-offs between theory and practice: fancy optimizers or deeper models don't always outperform simple baselines without good data and proper tuning.
- **2 more weeks**
  - Collect and integrate real user interaction data (e.g., clickstream or browsing logs) to make the cold-start problem more realistic.
- **Change one thing**
  - Spend more time on data pipeline and cleaning upfront; underestimate how much “data wrangling” would dominate the workload.

# Laura

- **Most Surprising Result or Finding**
  - RMSE didn't align with top-k recommendation quality - models with low RMSE often failed to rank relevant items effectively -> optimization objectives must be carefully chosen
- **Most useful lecture concept**
  - The SGD noise and preconditioning lectures helped us understand how to stabilize training with small batches, especially when using Adam in our Low rank and Two-Tower model.
- **Perspective change**
  - I shifted from trial-and-error tuning to a more systematic approach to guide choices on regularization, learning rates, and batch size for convergence.
- **2 more weeks**
  - We'd explore advanced optimizers (e.g., warm-up schedules, adaptive clipping) and test personalized regularization strategies to improve model generalization.
- **Change one thing**
  - We'd start by benchmarking non-matrix-factorization models (e.g., graph-based or transformer-based), to broaden the design space and better match metadata-rich recommendation scenarios.

# Matthew

- **Most Surprising Result or Finding**
  - Two towers algorithm isn't the best performing
- **Most useful lecture concept**
  - Transformer visualization website
  - How the K,Q,V matrices are used
- **Perspective on Optimization**
  - Lots of places can go wrong -> need to be careful and only change one thing at a time and understand why
- **2 more weeks**
  - Standardize the images to have better image embeddings
- **Change one thing**
  - Spend more time on generating data and making sure the format is consistent

Questions?